

Hermes Agent for TelemetryFlow Observability

Self-Improving AI Agent with Memory, Skills & Multi-Agent Teams

Platform: TelemetryFlow — Community Enterprise Observability Platform

Agent: Hermes — Self-Improving AI Agent Framework by Nous Research

Signals: Metrics, Logs, Traces, Exemplars

The Observability Challenge

Why Observability Needs AI

Modern distributed systems generate millions of telemetry signals per second — but the tools to make sense of them haven't kept pace with the complexity they create.

Challenge	Impact
Signal Flood	Thousands of alerts per hour — most are noise
Manual Correlation	Engineers switch between 4+ tools to trace root cause
Slow MTTR	Average incident resolution: 30–45 minutes
Knowledge Silos	Expert context lives in heads, not systems
No Automated Action	Investigation is manual; remediation requires human steps

You have all the data. You don't have an agent that can **read it all, correlate it, remember it, and act on it** — at machine speed.

TelemetryFlow Platform Overview

TelemetryFlow is an **enterprise-grade, 100% OTLP-compliant** observability platform that unifies four telemetry signals into a single pane of glass.

Signal	Storage	Capability
Metrics	ClickHouse	Time-series with interactive charts, zoom/pan
Logs	ClickHouse	Real-time streaming, full-text search, severity filtering
Traces	ClickHouse	Waterfall, flame graph, service dependency graph
Exemplars	ClickHouse	Metric-to-trace links for context-aware debug

- **Multi-Tenancy**: Region → Organization → Workspace → Tenant
- **33 Production Alert Rules** with fatigue prevention
- **6 Dashboard Templates** with 12+ widget types
- **Enterprise Security**: JWT, MFA, SSO, 5-Tier RBAC, API Keys
- **Database Monitoring**: Native collectors for 9 databases with QAN
- **TFO Agent**: Replaces Prometheus + KSM + node-exporter + FluentBit + cAdvisor
- **AI Intelligence**: MCP servers for Claude AI, TFQL natural language queries
- **TFO Collector v1.2.1**: OCB-Native Gateway with 4 custom TFO components

Introducing Hermes Agent

What is Hermes Agent?

A **self-improving AI agent framework** by Nous Research. The agent accumulates memory across sessions, writes and prunes its own reusable skills, and can be optimized offline via an evolutionary pipeline (GEPA). No other open-source agent ships all three in a single framework.

Property	Detail
License	Apache 2.0
Source	github.com/NousResearch/hermes-agent
Core	Single <code>AIAgent</code> class in <code>run_agent.py</code> — ReAct loop
Providers	9 LLM providers — Anthropic, OpenAI, OpenCode Go, OpenRouter, Google, Mistral, Groq, Ollama, LM Studio
Execution	6 backends — local, Docker, SSH, Modal, Daytona, Singularity
Turn Cap	90-turn hard cap per task (prevents runaway loops)
Memory	3-tier system — Persistent Markdown + SQLite FTS5 + External providers
Skills	687 skills across 18 categories (self-evolving + hub)
Multi-Agent	Profiles + delegation tool for parallel teams
Scheduling	Cron in plain English — jobs survive restarts

Architecture Overview

Everything flows through a single `AIAgent` class. CLI, Telegram gateway, batch runner, and IDE integrations are all entry points into the same core loop.

```
graph TD
    EP["<b>Entry Points</b><br>CLI · Telegram · Batch Runner · IDE"]
    CORE["<b>AIAgent Core</b> (run_agent.py)<br>ReAct loop: Build prompt → API call → Execute tools → Loop"]
    SOUL["SOUL.md<br><i>identity</i>"]
    MEM["Memory<br><i>3 tiers</i>"]
    SKILLS["Skills<br><i>procedural memory</i>"]
    EP --> CORE
    CORE --> SOUL
    CORE --> MEM
    CORE --> SKILLS
```

Key design decisions:

- **90-turn hard cap** per task — prevents runaway loops silently burning API credits. Subagents share the same budget.
- **6 execution backends** — local terminal, Docker, SSH, Modal, Daytona, Singularity. Same code, just a config change.
- **Model-agnostic** — a translation layer routes any provider through one of three API formats. Swap between

Supported LLM Providers

9 providers with a unified translation layer — swap models with a single config change.

Provider	Example models	Best for	Local?
Anthropic	<code>claude-sonnet-4-5</code> , <code>claude-opus-4-5</code>	Highest quality reasoning, long context	No
OpenAI	<code>gpt-4o</code> , <code>gpt-4o-mini</code> , <code>o3-mini</code>	Strong general capability, wide ecosystem	No
OpenCode Go	<code>glm-5.1</code> , <code>glm-5.0</code>	Cost-efficient coding agent pipeline	No
OpenRouter	Any model via unified endpoint	Access to 200+ models via one API key	No
Google	<code>gemini-2.5-pro</code> , <code>gemini-2.5-flash</code>	Multimodal tasks, large context window	No
Mistral	<code>mistral-large</code> , <code>mistral-small</code>	European data-residency, efficient models	No
Groq	<code>llama-3.3-70b</code> , <code>mistral-8x7b</code>	Fastest inference (LPU hardware)	No
Ollama	Any pulled model, e.g. <code>llama3.3</code>	Full offline / air-gapped operation	Yes
LM Studio	Any loaded model	Local GUI-managed models	Yes

OpenRouter tip: Acts as a single endpoint for hundreds of models. Set `provider: openrouter` once and switch models by name without juggling multiple API keys.

Identity Layer — SOUL.md

Before memory loads, before skills load — **identity loads first.**

Location: `~/.hermes/SOUL.md`

Position: Slot #1 in the system prompt, always.

SOUL.md defines the agent's personality, tone, communication style, and hard limits. It is hand-authored and static — you write it once, tweak over time. Without it, Hermes falls back to a built-in default identity.

Why It Matters for Observability

Everything the agent writes — memory entries, skill files, RCA responses — happens through the lens of this identity. SOUL.md is the fixed frame; memory and skills are the moving parts inside it.

Example for TelemetryFlow RCA Agent:

```
# SOUL.md
```

```
You are a pragmatic senior SRE with strong observability expertise.  
You optimize for signal over noise, root cause over symptoms,  
and actionable remediation over advisory commentary.  
You never fabricate a metric. You never guess at a root cause  
without evidence from ClickHouse.
```

Memory System — Three Tiers

Tier 1 — Persistent Markdown (Always in Context)

Two tiny files injected as a frozen snapshot at session start. Always present, zero retrieval cost.

File	Max Size	Purpose
<code>~/.hermes/memories/MEMORY.md</code>	2,200 chars	Agent notes: project conventions, tool quirks, lessons learned
<code>~/.hermes/memories/USER.md</code>	1,375 chars	Your profile: name, preferences, skill level, things to avoid

- New entries written mid-session persist to disk immediately but only appear in the system prompt on the **next session**.
- At ~80% capacity, the agent **consolidates** — merging related entries into denser, more information-packed versions.

Observability Example

MEMORY.md

- payments-api always OOMs on deploy (4x in 30 days)
- workspace_id is mandatory on all ClickHouse queries
- node-pool-3 has recurring memory pressure
- alert fatigue rule: suppress < medium severity during deploys

Tier 2 — Full-Text Session Search (SQLite)

Every conversation (CLI and Telegram) is stored in `~/hermes/state.db` with **FTS5 full-text search**. The agent can search weeks of past conversations on demand.

Tier	Trade-off
Tier 1	Always in context, tiny capacity
Tier 2	Unlimited capacity, requires active search + LLM summarisation

Critical facts live in Tier 1. **Everything else** is searchable on demand.

Observability Use Case

```
Agent: "What did we find when auth-service crashed last Tuesday?"  
→ Hermes searches state.db with FTS5  
→ Summarises the previous RCA  
→ Applies learnings to the current investigation
```

Tier 3 — External Memory Providers

Eight pluggable providers that run alongside built-in memory (never replacing it). Only one can be active at a time.

When active, Hermes automatically:

- **Prefetches** relevant memories before each turn
- **Syncs** conversation turns after each response
- **Extracts** memories on session end

For TelemetryFlow deployments, this enables **cross-cluster memory** — an agent in us-east-1 can recall investigation patterns from ap-southeast-1 through a shared memory provider.

Self-Evolving Skills

Skills: Procedural Memory

Memory handles **facts**. **Skills handle procedures**.

A skill is a Markdown file with YAML frontmatter — the agent's procedural memory for **how** it does things.

```
---
name: k8s-pod-debug
description: >
  Activate for crashing pods, CrashLoopBackOff,
  "why is my pod restarting", container failures.
version: 1.2.0
author: agent
platforms: [linux, macos]
---

## Procedure

1. Get pod status → check events → pull logs
2. Look for OOMKilled, ImagePullBackOff, config errors

## Pitfalls

- Forgetting --previous flag on restarted containers

## Verification
```

Self-Improvement Loop

Skill creation triggers **automatically** when:

- The agent completes a complex task (5+ tool calls)
- It hits errors or dead ends and finds the working path
- The user corrects its approach
- It discovers a non-trivial workflow

```
graph LR
  A[Encounter Problem] --> B[Solve via Trial & Error]
  B --> C[Save Successful Approach as SKILL.md]
  C --> D[Next Similar Problem]
  D --> E[Load Skill Instead of Rediscovering]
  E --> A
```

The loop: encounter problem → solve through trial and error → save successful approach as `SKILL.md` → next similar problem loads the skill instead of rediscovering from scratch.

The `skill_manage` tool supports: `create`, `patch`, `edit`, `delete`, `write_file`, `remove_file`.

| `patch` is the preferred action — targeted fix, token-efficient. Use `edit` only for full rewrites.

The Curator — Skill Garbage Collection

Without maintenance, agent-created skills pile up into dozens of narrow, overlapping playbooks that waste tokens and pollute the catalog.

Trigger: Inactivity check (not cron) — if 7 days have passed since last run AND the agent has been idle for 2+ hours, a background fork spins up with its own prompt cache, never touching the active conversation.

Two-Phase Operation

Phase 1 — Automatic Transitions (no LLM, deterministic):

- Unused for 30 days → `stale`
- Unused for 90 days → `archived`

Phase 2 — LLM Review (up to 8 iterations):

- A forked agent surveys all agent-authored skills
- Per-skill decision: keep, patch, consolidate, or archive

Safety Constraints

- **Never touches** bundled or hub-installed skills — only agent-authored ones
- **Never auto-deletes** — worst outcome is archival to `~/.hermes/skills/.archive/` (recoverable with one command; rollbacks are themselves reversible)

GEPA — Offline Skill Optimization

The in-agent learning loop has a known weakness: **the agent tends toward self-congratulation** — it almost always thinks it performed well, even when it didn't.

GEPA (**Genetic-Pareto Prompt Evolution**) is the solution. It lives in a **companion repository** (`NousResearch/hermes-agent-self-evolution`), operates fully offline, and was published as an **ICLR 2026 Oral paper** (MIT licensed).

```
graph TD
  A["Read current skill"] --> B["Generate evaluation dataset<br/><i>synthetic via Claude Opus · real SQLite history · hand-curated golden sets</i>"]
  B --> C["GEPA optimizer: read traces → identify failure points → generate candidate variants"]
  C --> D["LLM-as-judge scoring (rubrics, not binary pass/fail)"]
  D --> E["Constraint gates: full test suite passes 100%, skill stays under 15KB, caching compatibility preserved, semantic purpose doesn't drift"]
  E --> F["Best variant → PR against Hermes repo (never a direct commit)"]
```

Cost: ~\$2–10 per optimization run. No GPU required — everything runs through API calls.

Skip GEPA initially. Use it when you hit a wall and don't want to spend time and money on fine-tuning (RL/GRPO). It's a strong alternative to try before moving to full RL-based fine-tuning.

Skills Hub

Hermes maintains an official Skills Hub with **687 skills across 18 categories** — plus TelemetryFlow Hermes bundles **29 specialised skills** across **18 categories** covering all 20 TFO Platform modules:

Source	Count
Built-in (ships with agent)	87
Optional (enable on demand)	79
Anthropic (frontend-design, pdf, pptx, docx, mcp-builder, etc.)	16
LobeHub (community)	505

Add any GitHub repo as a custom tap:

```
hermes skills tap add yourname/your-skills-repo
hermes skills install yourname/your-skills-repo/<skill-name>
```

For TelemetryFlow, this means you can install observability-specific skills from the hub or create custom skill taps for your team's investigation patterns.

Multi-Agent Team for Observability

One Agent vs a Team

A single agent doing everything fills its context window fast — by step 4 it's operating on summaries of summaries and quality degrades. Separate concerns: each agent has one job, one context, one set of tools.

Hermes uses **profiles** — each with its own Telegram bot, SOUL.md, memory, and skills. They share nothing.

There are two distinct ways to use multiple agents. Both use Hermes profiles. The difference is what you're building:

Pattern	Agents	Use Case
Dev Pipeline	Supervisor → Coder → Reviewer → QC/Tester	Autonomous software delivery pipeline
Personal	Designer + Programmer + Researcher	Specialised assistants
Observability	Triage → Investigator → Reviewer → Remediator	AI-driven RCA & remediation pipeline

Observability Team Architecture

Adapted for TelemetryFlow: four specialised agents forming an autonomous incident response pipeline, powered by **40 plugin tools** covering all **20 TFO Platform modules**.

```
graph TD
    ALERT["Alert Fired<br/>(TFO Alert Engine / PagerDuty / Slack / K8s Events)"]
    TRIAGE["<b>Triage Agent</b><br/>glm-5.1 – classifies severity & routes"]
    INV["<b>Investigator</b><br/>claude-sonnet-4-5 + 40 plugin tools – metrics, logs, traces, exemplars, K8s, DB, LLM"]
    REV["<b>Reviewer</b><br/>glm-5.1 – separate context, read-only, zero investigation bias"]
    REM["<b>Remediator</b><br/>glm-5.1 + K8s tools – propose & gate remediation (human approval)"]
    HUMAN["<b>Human</b><br/>Approve / Reject / Manual review"]

    ALERT --> TRIAGE
    TRIAGE -->|"needs investigation?"| INV
    INV --> REV
    REV --> REM
    REM --> HUMAN
```

Steps 1–4 are fully autonomous. You only touch step 5.

Agent 1 — Triage (Hermes)

Detects new alerts, reads them, and decides: critical incident, known pattern, or noise? Does not investigate — it manages the queue and delegates.

SOUL.md excerpt:

```
You are an incident triage specialist. You classify alerts by severity
and route them to the appropriate agent. You never investigate – you
decide WHO should investigate. Known patterns from MEMORY.md get
auto-resolved. Only genuine anomalies get escalated.
```

Behavior:

- Reads alert payload from TelemetryFlow webhook
- Cross-references MEMORY.md for known patterns
- Classifies: critical / known / noise
- Delegates to Investigator only for genuine anomalies
- Auto-resolves known patterns (e.g., "payments-api OOM on deploy — known issue, rollback initiated")

Agent 2 — Investigator (Hermes + Plugin Tools)

The workhorse. Gets an alert, queries TelemetryFlow's four signals via dedicated plugin tools, correlates evidence, produces root cause hypothesis.

Skills auto-created after investigations:

```
---
name: payments-api-oom-rca
description: >
  Activate when payments-api shows OOM kill pattern,
  memory spike, or pod restart cascade.
version: 1.3.0
author: agent
---
```

Procedure

1. `query_metrics --signal metrics --service payments-api` → detect memory spike
2. `search_logs --severity ERROR` → find OOM kill messages
3. `list_traces --min-duration 500` → identify slow spans
4. `get_exemplars --metric memory usage` → link to specific traces
5. `query_correlations` → cross-signal evidence
6. Cross-reference with MEMORY.md: previous OOM incidents

Root Cause Pattern

- 512 MiB memory request insufficient after v2.4.1 deploy
- Previous occurrences: v2.3.0, v2.3.4, v2.4.0

Agent 3 — Reviewer (Subagent)

Runs in a **separate context with zero investigation bias** — it only sees the evidence and the hypothesis, not what the Investigator was thinking. Uses **read-only tools** to avoid bias toward defending its own conclusions.

Key design: A reviewer that can modify the investigation gets biased toward defending its own edits.

What the Reviewer checks:

- Does the evidence support the root cause hypothesis?
- Are there alternative explanations not considered?
- Is the proposed remediation safe and proportional?
- Does MEMORY.md contain conflicting historical data?

Agent 4 — Remediator (Hermes + K8s Tools)

The final gate before human review. Proposes concrete remediation actions, each requiring human approval.

Actions (gated):

Action	Tool	Gate Required
Scale deployment	<code>scale_deployment</code>	Human approval
Restart pod	<code>restart_pod</code>	Human approval
Rollback deploy	<code>rollback_deploy</code>	Human approval
Update alert rule	<code>update_alert</code>	Human approval
Escalate to human	<code>escalate</code>	Automatic (no gate)

Human-in-the-Loop notification (via Telegram):

- Alert summary + root cause
- Evidence links (metrics, logs, traces)
- Proposed remediation with risk assessment
- Reviewer verdict
- One-click: **Approve** / **Reject** / **Manual Review**

RCA & Postmortem — 3 Reporting Tools

After remediation, the pipeline generates structured incident reports with zero manual effort.

Tool	Output
<code>generate_rca_report</code>	Full Root Cause Analysis — 5W analysis, mermaid timeline, impact metrics, blast radius, action items
<code>generate_postmortem</code>	Comprehensive postmortem — timeline, lessons learned, remediation flow, appendix
<code>generate_rca_template</code>	Blank template for manual human review — document control, signature section

Jira & Trello Integration

`generate_rca_report` can create Jira tickets and Trello cards alongside the RCA report. Submission is **generate-only by default** — dual-gated by:

1. **Master switch env vars:** `JIRA_ENABLED=true` / `TRELLO_ENABLED=true` (off by default)
2. `--submit true` **flag**: explicit opt-in per invocation

Actions:

```
rca           → report only
jira / trello → report + ticket (generated, not submitted)
jira-submit / trello-submit / submit → force submit regardless of flags
```

Report Artifacts

Cost Per Incident

Agent	Model	Avg Cost
Triage	GLM 5.1 via OpenCode Go	~\$0.01
Investigator	Claude Sonnet 4.5 / GPT-4o	~\$0.05–0.15
Reviewer	GLM 5.1 via OpenCode Go	~\$0.03–0.08
Remediator	GLM 5.1 via OpenCode Go	~\$0.01–0.03
Total		~\$0.10–0.27 / incident

Use expensive models only for investigation (the complex reasoning task).

Triage, review, and remediation can use cheaper/faster models.

Original estimate with Claude Sonnet was ~\$0.39/ticket. GLM 5.1 may require more turns on complex bugs but costs ~3–4x less.

Live Investigation Walkthrough

Scenario: payments-api Latency Breach

3:47 AM — Alert fires: `payments-api` p95 latency breach (180ms → 640ms)

Step 1 — Triage Agent (Autonomous)

```
Triage reads alert payload.  
Cross-references MEMORY.md: "payments-api OOM pattern – 4x in 30 days"  
Severity: HIGH – genuine anomaly (different from past patterns: latency, not OOM)  
→ Delegates to Investigator
```

Step 2 — Investigator Agent (Autonomous)

```
Loads skill: payments-api-oom-rca (auto-created from past investigations)  
Queries TelemetryFlow via plugin tools:  
  query_metrics --signal metrics → memory spike from 512MiB to 890MiB  
  search_logs --severity ERROR → 512 OOM kill messages  
  list_traces --min-duration 500 → 89 slow spans on pod payments-api-7b8cf  
  get_exemplars --metric memory_usage → linked traces show allocation burst  
→ Root cause hypothesis: OOM after v2.4.1 deploy, memory request insufficient
```

Walkthrough (continued)

Step 3 — Reviewer Agent (Autonomous, Separate Context)

Reviewer sees ONLY the evidence and hypothesis – not the Investigator's thought process.

Read-only verification:

- ✓ Metrics confirm memory spike
- ✓ Logs confirm OOM kill
- ✓ Traces confirm latency correlation
- ✓ MEMORY.md confirms recurring pattern

⚠ Alternative: could be a memory leak rather than insufficient request

→ Verdict: root cause likely correct, recommend memory profiling as follow-up

Step 4 — Remediator Agent (Gated)

Proposes remediation:

- Roll back to v2.4.0 [REQUEST APPROVAL]
- Or raise memory to 1 GiB [REQUEST APPROVAL]
- Long-term: add memory profiling to CI pipeline

Sends Telegram notification to on-call engineer.

Step 5 — Human Decision

Engineer receives Telegram notification:

MTTR: The Numbers

Stage	MTTR	Improvement
Manual Investigation	~40 min	Baseline
Single AI Agent	~47 sec	50× faster
Hermes Team (4 agents)	~23 sec	100× faster

Why It's Faster

- **Specialised agents:** Each agent has one job, one context, one set of tools — no context window pollution
- **Skill loading:** Past investigation procedures loaded instantly instead of rediscovered
- **Memory recall:** Known patterns auto-resolved at triage — only genuine anomalies reach investigation
- **Action gate:** Human approves while the pipeline is ready to execute immediately

The engineer reviews an AI-produced hypothesis with evidence AND a one-click remediation — instead of assembling either from scratch.

Deployment & Configuration

Install & Configure Hermes

Step 1 — Install

```
curl -fsSL https://raw.githubusercontent.com/NousResearch/hermes-agent/main/scripts/install.sh | bash
source ~/.bashrc

# Verify – all checks should be green
hermes doctor

# Auto-repair if any issues
hermes doctor --fix

# Launch interactive chat
hermes

# Type /help to see all slash commands
```

Step 2 — Configure LLM Provider

```
# Interactive setup wizard (recommended first time)
hermes setup model

# Or set directly
hermes config set model.default "claude-sonnet-4-5"
hermes config set model.provider "anthropic"
```

Configure LLM Provider (API Keys)

Edit `~/hermes/.env` — add only the key for your chosen provider:

```
# Anthropic
ANTHROPIC_API_KEY=your_key_here

# OpenAI
OPENAI_API_KEY=your_key_here

# OpenCode Go (GLM)
OPENCODE_API_KEY=your_key_here

# OpenRouter (covers many providers via one key)
OPENROUTER_API_KEY=your_key_here
```

Switching Models Mid-Project

```
# Check current model
hermes config get model.default

# Switch to a different provider/model on the fly
hermes config set model.default "gpt-4o-mini"
hermes config set model.provider "openai"
```

TelemetryFlow Plugin Tools — 40 Tools, 20 Modules

All tools are Python stdlib only (zero dependencies). Each tool communicates with TFO Platform via REST API — no direct ClickHouse connections.

Category	Tools
Core Telemetry (5)	<code>query_metrics</code> , <code>search_logs</code> , <code>list_traces</code> , <code>get_exemplars</code> , <code>query_correlations</code>
Monitoring (8)	<code>check_k8s</code> , <code>check_infra</code> , <code>check_uptime</code> , <code>check_vm</code> , <code>check_agent</code> , <code>check_service_map</code> , <code>check_network_map</code> , <code>check_db_monitoring</code>
AI & LLM (7)	<code>chat_with_context</code> , <code>stream_chat</code> , <code>manage_conversation</code> , <code>generate_insight</code> , <code>query_llm_usage</code> , <code>manage_provider</code> , <code>query_ai_intelligence</code>
Platform (8)	<code>query_platform</code> , <code>query_account</code> , <code>query_audit</code> , <code>query_subscription</code> , <code>manage_dashboards</code> , <code>manage_alerts</code> , <code>manage_reports</code> , <code>manage_data_masking</code>
Infrastructure (5)	<code>manage_retention</code> , <code>manage_tenancy</code> , <code>manage_iam</code> , <code>manage_sso</code> , <code>query_tfql</code>
Remediation (4) ⚠️	<code>scale_deployment</code> , <code>restart_pod</code> , <code>rollback_deploy</code> , <code>update_alert</code> — all require human approval
RCA & Postmortem (3)	<code>generate_rca_report</code> , <code>generate_postmortem</code> , <code>generate_rca_template</code> — 5W analysis, mermaid diagrams, Jira/Trello integration

74 ContextType values from TFO's ContextCollector, **15 LLM provider types**, **5 insight types** — full coverage of the TFO LLM module.

Connect Telegram Gateway

Prerequisites:

- Bot token from @BotFather
- Your chat ID → message your bot, then visit `https://api.telegram.org/bot<TOKEN>/getUpdates`

```
hermes gateway setup    # select Telegram, enter token + chat ID

hermes gateway start
hermes gateway status
```

Send a message to your bot on Telegram, then run `/sethome`.

Enable Required Tools

```
hermes tools enable terminal    # run shell commands (ClickHouse queries)
hermes tools enable cronjob     # schedule investigations
hermes tools enable delegation  # spawn subagents (multi-agent team)
hermes tools enable web         # hit TelemetryFlow API and external APIs

hermes tools list               # verify
```

Run `/reset` after enabling — changes apply to new sessions only.

Create Observability Agent Profiles

Each agent in the team gets its own Hermes profile with dedicated SOUL.md, memory, and skills.

```
# Create profiles for the observability team
hermes profile create triage      --clone
hermes profile create investigator --clone
hermes profile create reviewer   --clone
hermes profile create remediator --clone

hermes profile list
```

`--clone` copies your default profile's config and `.env` as a starting point.

Configure Each Profile

```
# Triage – cheap/fast model
hermes -p triage config set model.default "glm-5.1"
hermes -p triage config set model.provider "opencode-go"

# Investigator – powerful model for complex reasoning
hermes -p investigator config set model.default "claude-sonnet-4-5"
hermes -p investigator config set model.provider "anthropic"

# Reviewer – cheap model for read-only verification
```

Give Each Profile Its Own Telegram Bot

Telegram allows only one connection per token. Run `/newbot` four times with BotFather, then:

```
hermes -p triage      gateway setup
hermes -p investigator gateway setup
hermes -p reviewer    gateway setup
hermes -p remediator  gateway setup
```

Write SOUL.md per Profile

Edit each profile's `SOUL.md` to make the agents genuinely different:

- **Triage** (`~/.hermes/profiles/triage/SOUL.md`): "You are an incident triage specialist. You classify alerts by severity and route them. You never investigate — you decide WHO should investigate."
- **Investigator** (`~/.hermes/profiles/investigator/SOUL.md`): "You are a senior SRE investigator. You query ClickHouse for evidence. You never guess at a root cause without data."
- **Reviewer** (`~/.hermes/profiles/reviewer/SOUL.md`): "You are an independent reviewer. You only see evidence and hypotheses. You check for bias and missed alternatives."
- **Remediator** (`~/.hermes/profiles/remediator/SOUL.md`): "You are a remediation specialist. You propose gated actions. Every write operation requires human approval."

Connect TelemetryFlow Alert Webhooks

Wire TelemetryFlow alert rules to the Triage agent's Telegram gateway.

Step 1 — Set Up Telegram Gateway

```
hermes -p triage gateway setup    # select Telegram, enter token + chat ID
hermes -p triage gateway start
```

Step 2 — Configure TelemetryFlow Webhook

In TelemetryFlow Alert Rules, set the notification webhook to forward alert payloads to the Triage agent's endpoint.

Step 3 — Schedule Periodic Health Checks

```
# Plain-English cron – no cron expressions needed
hermes -p investigator cron add "every 15m" "Check TelemetryFlow ClickHouse for anomaly spikes in metrics_1m"
```

Step 4 — Verify the Pipeline

Send a test alert through TelemetryFlow. Hermes Triage should classify it, and if genuine, the full pipeline should execute through to the human approval notification.

Scheduling — Cron in Plain English

The Hermes scheduler ticks every 60 seconds, runs due jobs in isolated agent sessions, and delivers output to your messaging platform. Jobs survive restarts. They live in `~/.hermes/cron/jobs.json`.

You don't write cron expressions — you describe what you want in plain English and Hermes converts it.

Cron Syntax Reference

Pattern	Example
One-shot delay	<code>/cron add 30m "Remind me to check the build"</code>
Recurring interval	<code>/cron add "every 2h" "Check server status"</code>
Standard cron expression	<code>/cron add "0 9 * * 1-5" "Weekday 9am task"</code>
Skill attachment	<code>/cron add "every 1h" "Summarize feed items" --skill blogwatcher</code>

Chain jobs: one cron's output becomes the next's input via a `context_from` flag. Useful for research → writing pipelines.

Security Model

Layer	Mechanism	Prevents
Credentials	<code>~/.hermes/.env</code> — secrets only	Passwords in code / config.yaml
DB Roles	Read-only ClickHouse user for Hermes	Unbounded data mutation
Tool Permissions	40 plugin tools, 4 gated ⚠️	Unauthorized write access
Action Gates	Human-in-the-loop for write operations	Autonomous mutation
Tenant Isolation	Mandatory workspace_id on all queries	Cross-tenant data leakage
Reviewer Bias	Separate context, read-only tools	Investigation bias defense
Turn Cap	90-turn hard cap per task	Runaway loops / credit burn
Sandbox	Container-level network policy	Lateral movement
Curator Safety	Snapshot + pin + archive (never delete)	Skill loss
Cyber Defense	Threat-informed triage + security review on every investigation	Security blind spots

Air-Gapped Deployment

For environments requiring zero external egress, use **Ollama** as the LLM provider. Or deploy with **Docker** using the provided Dockerfile and docker-compose.yaml.

Docker Deployment

```
# Build and start with TFO Platform stack
./run-container.sh -b --up --profile core

# Or full stack (backend + frontend + collector + agent)
./run-container.sh -b --up --profile all
```

```
graph LR
    HERMES["tfo-hermes<br/>40 Plugin Tools"] --> API["TFO Backend<br/>:3000"]
    API --> PG["PostgreSQL"]
    API --> CH["ClickHouse"]
    API --> RD["Redis"]
    API --> NATS["NATS"]

    COLLECTOR["tfo-collector<br/>:4317/:4318"] --> API
    AGENT["tfo-agent<br/>Host Metrics"] --> COLLECTOR
    VIZ["tfo-viz<br/>:8080"] --> API
```

TelemetryFlow + Hermes: The Full Picture

TelemetryFlow Feature Coverage

How Hermes plugin tools map to TelemetryFlow's complete feature set — all 20 TFO Platform modules covered.

TelemetryFlow Feature	Plugin Tool	How It Works
Metrics Explorer	<code>query_metrics</code>	Query metrics_1m/5m/1h via TFO API
Log Viewer	<code>search_logs</code>	FTS search on otel_logs with severity filters
Distributed Tracing	<code>list_traces</code>	Query otel_traces for span waterfall analysis
Exemplars	<code>get_exemplars</code>	Join metrics and traces via exemplar links
Alert Management	<code>manage_alerts</code>	CRUD alert rules, verify thresholds
Dashboard Builder	<code>manage_dashboards</code>	Create/read/update dashboards and widgets
Uptime Monitoring	<code>check_uptime</code>	Monitors, status pages, SSL, incidents
VM Monitoring	<code>check_vm</code>	Virtual machine inventory and metrics
Agent Monitoring	<code>check_agent</code>	TFO Agent health, heartbeat, metrics
Kubernetes	<code>check_k8s</code>	Clusters, nodes, pods, deployments
Service Map	<code>check_service_map</code>	Service topology, dependencies, health
Network Map	<code>check_network_map</code>	Network topology, connections, flow
Database Monitoring	<code>check_db_monitoring</code>	16 database types + QAN

File System Layout

```
~/hermes/
├── config.yaml      # Model choice, terminal backend, tool enablement
├── .env            # API keys and secrets (never in config.yaml)
├── auth.json       # OAuth provider credentials
├── SOUL.md         # Agent identity (slot #1 in system prompt)
├── memories/
│   ├── MEMORY.md   # Persistent agent facts (2,200 char max)
│   └── USER.md     # User profile (1,375 char max)
├── skills/         # All skills (bundled, hub, agent-created)
│   ├── monitoring/ # 8 monitoring skills (uptime, vm, agent, k8s, service-map, network-map, ...)
│   ├── observability/ # 9 observability skills
│   │   ├── k8s-pod-debug/
│   │   │   └── SKILL.md # Auto-created investigation procedure
│   │   ├── payments-api-oom-rca/
│   │   │   └── SKILL.md
│   │   └── clickhouse-query-patterns/
│   │       └── SKILL.md
│   ├── alerting/   # alert-management
│   ├── dashboard/  # dashboard-management
│   ├── reporting/  # report-automation
│   ├── database-monitoring/ # slow-query-detection, qan-analysis
│   └── .hub/       # Skills Hub state (687 skills, 18 categories)
├── profiles/       # Multi-agent team profiles
│   ├── triage/     # Triage agent config + SOUL.md + memory
│   ├── investigator/ # Investigator agent config + SOUL.md + memory
│   ├── reviewer/   # Reviewer agent config + SOUL.md + memory
│   └── remediator/  # Remediator agent config + SOUL.md + memory
├── sessions/       # Per-platform session metadata
├── state.db        # SQLite + FTSS session store
├── cron/
│   ├── jobs.json   # Scheduled investigation jobs
│   └── output/     # Cron run outputs
├── plugins/        # TelemetryFlow plugin (40 tools)
│   └── telemetryflow/
│       ├── plugin.yaml # v1.2.0 - 40 tools, 3 env vars
│       └── tools/      # Python stdlib only (zero deps)
│           ├── _shared.py # API helpers, 74 ContextTypes, 15 ProviderTypes
│           ├── query_metrics.py
│           ├── search_logs.py
│           ├── ...      # 40 tools total covering all 20 TFO modules
│           ├── scale_deployment.py # ⚠ requires_approval
│           ├── restart_pod.py      # ⚠ requires_approval
│           ├── rollback_deploy.py  # ⚠ requires_approval
│           └── update_alert.py      # ⚠ requires_approval
├── hooks/          # Lifecycle hooks
├── skins/          # CLI themes
└── logs/           # agent.log, gateway.log, errors.log
```

Troubleshooting

Symptom	Fix
Gateway not responding	<code>grep -i error ~/.hermes/logs/gateway.log tail -20</code> → <code>hermes gateway restart</code>
Tool not available	<code>hermes tools list</code> → enable missing tools → <code>/reset</code>
Model not responding	<code>hermes doctor</code> → check <code>~/.hermes/.env</code> credentials
Gateway dies on SSH disconnect	<code>sudo loginctl enable-linger \$USER</code>
Skills accumulating / stale catalog	<code>hermes curator run</code> to trigger manual Curator pass
Skill overwritten with worse version	Check <code>~/.hermes/skills/.archive/</code> — rollback with one command
GEPA optimization cost too high	Start with real SQLite history as eval set; avoid full synthetic generation

Key Takeaways

1. **TelemetryFlow already has the data** — Metrics, Logs, Traces, Exemplars, all in ClickHouse, all OTLP-compliant. No pipeline changes needed.
2. **Hermes is a self-improving agent** — It accumulates memory across sessions, writes and prunes its own reusable skills, and can be optimized offline via GEPA. No other open-source agent ships all three.
3. **The integration is plugin tools + API** — Hermes queries TelemetryFlow through 40 dedicated plugin tools covering all 20 TFO Platform modules. No direct ClickHouse connections — everything goes through the authenticated REST API.
4. **Multi-agent teams separate concerns** — Triage, Investigation, Review, and Remediation each get their own profile, context, and model. No context window pollution.
5. **Security is layered** — Read-only DB roles, tool permissions, human-approval gates, 90-turn cap, reviewer bias protection, and skill curator safety.
6. **Skills compound over time** — Every successful investigation becomes a reusable procedure. The agent gets better at RCA with every incident, automatically.
7. **Scheduling is built-in** — Plain-English cron for periodic health checks, chained jobs for research → remediation pipelines. Jobs survive restarts.

Questions?

TelemetryFlow + Hermes: Intelligent Observability at Machine Speed

- The observability data is already there.
- The self-improving agent is already there.
- The integration is 40 plugin tools + TFO REST API.
- The evolution is one webhook switch away.

Dwi Fahni Denni · TelemetryFlow · 2026

Resources:

- github.com/NousResearch/hermes-agent
- github.com/telemetryflow
- telemetryflow.id